

NVIDIA Base Command Manager - Detailed Procedures & Solutions for NCP-AIOPS

1. Installation & Initial Configuration

Pre-Installation Requirements

```
bash
```

```
# System Requirements Check
```

- Head Node: 64GB RAM minimum, 1TB storage
- Compute Nodes: GPU-enabled systems with CUDA drivers
- Network: InfiniBand or high-speed Ethernet
- OS: RHEL 8.x, Ubuntu 20.04+ or SLES 15

```
# Network Prerequisites
```

- DHCP server configuration
- DNS resolution setup
- NTP synchronization
- Firewall rules for cluster communication

Installation Process

```
bash
```

1. Download and extract BCM installer

`wget https://nvidia.com/base-command-manager-installer.tar.gz`

`tar -xzf base-command-manager-installer.tar.gz`

2. Run installation wizard

`./install --wizard`

Follow prompts for:

- Network configuration

- Storage setup

- License activation

- Initial admin user creation

3. Verify installation

`cm-cluster-status`

`cm-node-list`

Post-Installation Configuration

`bash`

Configure cluster settings

`cm-cluster-configure --dns-domain cluster.local`

`cm-cluster-configure --timezone UTC`

Setup shared storage

`cm-storage-configure --type nfs --path /shared`

Initialize GPU monitoring

`cm-gpu-initialize --all-nodes`

2. Cluster Provisioning & Node Management

Node Provisioning Process

```
bash
```

```
# 1. Add compute nodes to cluster
```

```
cm-node-add --hostname compute-01 --mac 00:11:22:33:44:55
```

```
cm-node-add --hostname compute-02 --mac 00:11:22:33:44:56
```

```
# 2. Assign software images
```

```
cm-image-assign --node compute-01 --image ubuntu-20.04-gpu
```

```
cm-image-assign --node compute-02 --image ubuntu-20.04-gpu
```

```
# 3. Boot nodes
```

```
cm-node-boot --node compute-01,compute-02
```

```
# 4. Monitor provisioning status
```

```
cm-node-status --verbose
```

Software Image Management

```
bash
```

```
# Create custom software image
cm-image-create --name "ai-workload-v1" --base ubuntu-20.04
cm-chroot ai-workload-v1

# Inside chroot environment:
apt update && apt install -y nvidia-docker2 kubernetes-node
pip install torch tensorflow rapids-cudf
exit

# Finalize and deploy image
cm-image-finalize ai-workload-v1
cm-image-deploy --image ai-workload-v1 --nodes compute-group-gpu
```

Hardware Discovery & Configuration

```
bash

# GPU discovery and configuration
cm-gpu-discover --all-nodes
cm-gpu-configure --enable-persistence-mode
cm-gpu-configure --power-limit 250W

# Network interface configuration
cm-network-configure --interface ib0 --ip-range 192.168.100.0/24
cm-network-configure --interface eth0 --dhcp

# Storage configuration
cm-storage-mount --type lustre --server storage-server:/lustre --mountpoint /scratch
```

3. Workload Management & Job Scheduling

Kubernetes Configuration

bash

Enable Kubernetes orchestration

cm-kubernetes-enable --version 1.25

cm-gpu-operator-install --version 23.3.2

Configure GPU scheduling

kubect! apply -f - <<EOF

apiVersion: v1

kind: ResourceQuota

metadata:

name: gpu-quota

spec:

hard:

nvidia.com/gpu: "8"

requests.memory: "64Gi"

requests.cpu: "32"

EOF

Deploy GPU-enabled workload

kubect! apply -f - <<EOF

apiVersion: v1

kind: Pod

metadata:

name: gpu-pod

spec:

containers:

- name: gpu-container

image: nvidia/cuda:11.8-runtime-ubuntu20.04

resources:

limits:

nvidia.com/gpu: 1

```
command: ["/bin/bash", "-c", "nvidia-smi && sleep 3600"]
EOF
```

Slurm Integration

```
bash

# Configure Slurm workload manager
cm-slurm-configure --partition gpu --nodes compute[01-04]
cm-slurm-configure --partition cpu --nodes compute[05-08]

# Create Slurm configuration
cat > /etc/slurm/slurm.conf <<EOF
ClusterName=ai-cluster
ControlMachine=head-node
SlurmUser=slurm

# GPU Partition
PartitionName=gpu Nodes=compute[01-04] Default=YES MaxTime=24:00:00 State=UP
NodeName=compute[01-04] CPUs=32 RealMemory=128000 Gres=gpu:4 State=UNKNOWN

# CPU Partition
PartitionName=cpu Nodes=compute[05-08] MaxTime=48:00:00 State=UP
NodeName=compute[05-08] CPUs=64 RealMemory=256000 State=UNKNOWN
EOF

# Submit GPU job example
sbatch --partition=gpu --gres=gpu:2 --mem=32G gpu-training.sh
```

Job Submission Examples

```
bash
```

```
# AI Training Job Script (gpu-training.sh)
#!/bin/bash
#SBATCH --job-name=ai-training
#SBATCH --partition=gpu
#SBATCH --gres=gpu:4
#SBATCH --mem=64G
#SBATCH --time=12:00:00

module load cuda/11.8
module load python/3.9

python train_model.py --gpus 4 --batch-size 128 --epochs 100

# HPC Simulation Job
#!/bin/bash
#SBATCH --job-name=hpc-sim
#SBATCH --partition=cpu
#SBATCH --ntasks=256
#SBATCH --mem-per-cpu=4G
#SBATCH --time=24:00:00

mpirun -np 256 simulation.exe --input config.xml
```

4. Monitoring & Performance Optimization

Real-Time Monitoring Setup

```
bash
```



```
# Enable comprehensive monitoring
cm-monitoring-enable --components dcgm,prometheus,grafana
cm-monitoring-configure --retention-days 30

# DCGM (Data Center GPU Manager) configuration
dcgmi discovery -l # List all GPUs
dcgmi stats -g 0 -e # Enable stats collection
dcgmi health -g 0 # Check GPU health

# Custom monitoring queries
dcgmi stats -g 0 --verbose # Detailed GPU metrics
dcgmi dmon -e 155,150,203,204 # Monitor specific metrics
```

Performance Metrics Collection

```
bash

# System performance monitoring
cm-perf-monitor --interval 60 --metrics cpu,memory,network,gpu
cm-perf-alert --metric gpu-temp --threshold 85 --action email

# Job performance analysis
cm-job-analyze --job-id 12345 --metrics gpu-util,memory-usage
cm-job-compare --job-ids 12345,12346 --output report.html

# Network performance testing
cm-network-test --test bandwidth --nodes compute[01-04]
cm-network-test --test latency --protocol rdma
```

Performance Optimization Techniques

```
bash
```

```
# GPU optimization
```

```
nvidia-smi -pm 1 # Enable persistence mode
```

```
nvidia-smi -pl 300 # Set power limit
```

```
nvidia-smi --auto-boost-default=DISABLED # Disable auto-boost
```

```
# Memory optimization
```

```
echo 'vm.swappiness=1' >> /etc/sysctl.conf
```

```
echo 'vm.zone_reclaim_mode=1' >> /etc/sysctl.conf
```

```
# Network optimization for AI workloads
```

```
echo 'net.core.rmem_max=134217728' >> /etc/sysctl.conf
```

```
echo 'net.core.wmem_max=134217728' >> /etc/sysctl.conf
```

```
echo 'net.ipv4.tcp_congestion_control=bbr' >> /etc/sysctl.conf
```

5. Troubleshooting Procedures

Common Node Issues & Solutions

Node Boot Failures

```
bash
```

Diagnosis steps

```
cm-node-status compute-01 --detailed
```

```
cm-log-view --node compute-01 --service dhcp
```

```
cm-log-view --node compute-01 --service pxe
```

Common fixes

1. MAC address mismatch

```
cm-node-edit compute-01 --mac-correct 00:11:22:33:44:57
```

2. Network boot issues

```
cm-pxe-regenerate --node compute-01
```

```
cm-dhcp-restart
```

3. Image corruption

```
cm-image-verify ubuntu-20.04-gpu
```

```
cm-image-rebuild ubuntu-20.04-gpu --force
```

GPU Recognition Issues

```
bash
```

GPU troubleshooting

nvidia-smi *# Check if GPUs are visible*

lspci | grep -i nvidia *# Verify hardware detection*

Driver issues

cm-driver-update --component nvidia --nodes compute-01

cm-node-reboot compute-01

CUDA toolkit verification

nvcc --version

/usr/local/cuda/samples/1_Uutilities/deviceQuery/deviceQuery

Workload Management Issues

Job Scheduling Problems

bash

```
# Slurm troubleshooting
sinfo -R # Check node reasons
squeue --start # Show job start times
sacct -j 12345 --format=JobID,State,ExitCode,DerivedExitCode
```

```
# Common fixes
```

```
# 1. Node drain issues
```

```
scontrol update NodeName=compute-01 State=RESUME
```

```
# 2. Resource allocation problems
```

```
scontrol show job 12345
```

```
scontrol update JobId=12345 TimeLimit=48:00:00
```

```
# 3. GPU allocation issues
```

```
scontrol show partition gpu
```

```
scontrol update PartitionName=gpu MaxTime=24:00:00
```

Kubernetes Issues

```
bash
```

Pod troubleshooting

```
kubectl describe pod gpu-pod
```

```
kubectl logs gpu-pod
```

```
kubectl get events --sort-by='.lastTimestamp'
```

GPU operator issues

```
kubectl get pods -n gpu-operator-resources
```

```
kubectl logs -n gpu-operator-resources gpu-operator-xxx
```

Node issues

```
kubectl describe node compute-01
```

```
kubectl get nodes -o wide
```

Performance Troubleshooting

GPU Underutilization

```
bash
```

Diagnosis

`nvidia-smi dmon -s pucvmet -i 1` *# Monitor GPU metrics*

`dcgmi dmon -e 150,155,203,204` *# DCGM monitoring*

Common causes and fixes

1. CPU bottleneck

`top -p $(pgrep python)`

`taskset -cp 0-31 $(pgrep python)` *# Pin to specific CPUs*

2. I/O bottleneck

`iotop -p $(pgrep python)`

Move data to faster storage or increase buffer sizes

3. Memory bandwidth issues

`numactl --hardware`

`numactl --cpubind=0 --membind=0 python train.py`

Network Performance Issues

`bash`

InfiniBand troubleshooting

`ibstat` *# Check IB port status*

`ibdiagnet` *# Network topology and health check*

`perfquery` *# Performance counters*

Ethernet troubleshooting

`ethtool eth0` *# Check link status and settings*

`iperf3 -s` *# Server mode for bandwidth testing*

`iperf3 -c server-ip -t 60` *# Client bandwidth test*

6. Security & Access Management

User Authentication Setup

```
bash

# LDAP integration
cm-auth-configure --type ldap --server ldap.company.com
cm-auth-configure --base-dn "dc=company,dc=com"
cm-auth-configure --bind-user "cn=admin,dc=company,dc=com"

# OIDC configuration
cm-auth-configure --type oidc --provider https://auth.company.com
cm-auth-configure --client-id cluster-auth --client-secret xxx

# Local user management
cm-user-add --username scientist1 --groups gpu-users,data-scientists
cm-user-quota --username scientist1 --gpu-hours 100 --storage 1TB
```

Role-Based Access Control

```
bash
```


Create custom roles

cm-role-create --name "ml-engineer" --permissions job-submit,job-monitor

cm-role-create --name "admin" --permissions cluster-admin,user-admin

Assign roles to users

cm-user-role-assign --username scientist1 --role ml-engineer

cm-group-role-assign --group data-team --role ml-engineer

Resource access control

cm-partition-access --partition gpu --groups gpu-users

cm-storage-access --path /datasets --groups data-scientists --mode ro

Security Hardening

bash

Network security

cm-firewall-enable --strict

cm-firewall-rule --service ssh --port 22 --source admin-network

cm-firewall-rule --service slurm --port 6817-6818 --source cluster-network

Certificate management

cm-cert-generate --ca cluster-ca --validity 365

cm-cert-deploy --service kubernetes --cert cluster-cert

cm-cert-renew --auto-enable --days-before 30

Audit logging

cm-audit-enable --events user-login,job-submit,admin-actions

cm-audit-configure --retention 90 --export syslog

7. Cloud & Hybrid Integration

Cloud Burst Configuration

```
bash

# AWS integration
cm-cloud-configure --provider aws --region us-west-2
cm-cloud-credentials --access-key AKIAXXXXXX --secret-key xxxxx
cm-cloud-template --instance-type p3.8xlarge --image ami-xxxxx

# Auto-scaling rules
cm-autoscale-enable --min-nodes 0 --max-nodes 10
cm-autoscale-rule --queue-length 5 --scale-up 2
cm-autoscale-rule --idle-time 600 --scale-down 1

# Burst job submission
sbatch --partition=cloud --constraint=aws gpu-training.sh
```

Multi-Cloud Management

```
bash

# Multiple cloud providers
cm-cloud-add --name azure --provider azure --region eastus
cm-cloud-add --name gcp --provider gcp --region us-central1

# Workload distribution
cm-scheduler-policy --prefer on-premise --fallback cloud
cm-cost-optimize --provider cheapest --constraint gpu-memory=32GB
```

8. Backup & Disaster Recovery

Configuration Backup

```
bash

# Automated backup configuration
cm-backup-configure --schedule daily --retention 30
cm-backup-include --path /etc/base-command-manager
cm-backup-include --path /home --exclude "*.tmp;*.cache"

# Manual backup
cm-backup-create --name "pre-upgrade-$(date +%Y%m%d)"
cm-backup-verify --name pre-upgrade-20231201

# Restore procedures
cm-backup-restore --name pre-upgrade-20231201 --target /tmp/restore
cm-cluster-restore --from /tmp/restore --components config,users,jobs
```

Disaster Recovery Procedures

```
bash
```

```
# Head node failure recovery
# 1. Install BCM on replacement hardware
# 2. Restore from backup
cm-cluster-restore --full --backup latest

# 3. Update DNS and network configuration
cm-network-update --head-node-ip 192.168.1.100

# 4. Restart cluster services
cm-service-restart --all
cm-node-wake --all-compute-nodes
```

9. Capacity Planning & Scaling

Resource Utilization Analysis

```
bash

# Historical usage analysis
cm-usage-report --period monthly --format csv
cm-gpu-utilization --nodes all --timeframe "last 30 days"
cm-job-efficiency --users all --metrics gpu-hours,cpu-hours

# Capacity forecasting
cm-forecast --growth-rate 20% --horizon 12-months
cm-bottleneck-analysis --identify compute,storage,network
```

Scaling Procedures

```
bash
```

Adding new compute nodes

```
cm-node-add-batch --hostnames compute[09-16] --image ai-workload-v2
```

```
cm-partition-expand --partition gpu --nodes compute[09-12]
```

```
cm-partition-expand --partition cpu --nodes compute[13-16]
```

Storage scaling

```
cm-storage-expand --filesystem /scratch --capacity +100TB
```

```
cm-storage-rebalance --filesystem /scratch
```

Network scaling

```
cm-network-upgrade --interface ib0 --speed 200Gb
```

```
cm-network-add --interface ib1 --topology fat-tree
```

10. Exam-Specific Scenarios

Scenario 1: GPU Allocation Issue

Problem: Jobs requesting GPUs remain in pending state

```
bash
```

Diagnosis

```
squeue | grep PD # Check pending jobs
```

```
sinfo -o "%P %a %I %D %T %N" # Partition status
```

```
scontrol show partition gpu
```

Solution steps

1. Check GPU availability: `sinfo -o "%n %G"`
2. Verify GRES configuration: `scontrol show node compute-01`
3. Update Slurm configuration if needed
4. Restart slurm services: `systemctl restart slurmd`

Scenario 2: Node Performance Degradation

Problem: Compute node showing poor performance

```
bash
```

```
# Systematic diagnosis
```

1. Check system resources: `top`, `free -h`, `df -h`
2. Monitor GPU health: `nvidia-smi`, `dcgm health -g 0`
3. Check network: `ibstat`, `ethtool eth0`
4. Review system logs: `journalctl -u kubelet -f`

```
# Performance optimization
```

1. Update drivers: `cm-driver-update --node compute-01`
2. Optimize power settings: `nvidia-smi -pl 300`
3. Check thermal throttling: `nvidia-smi -q -d TEMPERATURE`

Scenario 3: Kubernetes Pod Scheduling Issues

Problem: Pods stuck in pending state despite available resources

```
bash
```

```
# Troubleshooting workflow
```

```
kubectl describe pod stuck-pod # Check events
```

```
kubectl get nodes -o wide # Node status
```

```
kubectl describe node compute-01 # Resource allocation
```

```
# Common fixes
```

1. Taint removal: `kubectl taint nodes compute-01 key=value:NoSchedule-`
2. Label addition: `kubectl label nodes compute-01 gpu=true`
3. Resource quota adjustment: `kubectl edit resourcequota gpu-quota`

This detailed guide provides the practical, hands-on procedures and troubleshooting steps essential for the NCP-AIOPS certification exam. Each section includes real commands, configuration examples, and step-by-step problem resolution procedures that system administrators encounter in production environments.